

PrePro **Auto**

Complete Project & User Guide

v1.0.0 • 2026 • Author: Shivanshu Pandey

```
pip install prepro-auto
```

278+	5	5	18
Tests passing	AI Providers	Cleaning Stages	Transform Presets

RAM-aware • Reproducible & Auditable • Leakage-Safe Model Training

Section 1 — Project Deep-Dive	Section 2 — Complete User Guide
Architecture, methods overview, dataset limits, and design philosophy.	Step-by-step: notebook interface, web UI, every function explained, all outputs described.

Contents

1	What is PrePro Auto?	2
2	Project Goal & Design Philosophy	2
3	Target Users	3
4	System Architecture	3
5	End-to-End Workflow	5
6	How It Works (Methods at a Glance)	17
7	Dataset Size Limits	18
8	Installation	19
9	Data Privacy	19
10	Notebook Interface — All Functions	20
11	Web UI — Step-by-Step Guide	22
12	All Outputs Explained	24
13	AI Providers Setup	24
14	Troubleshooting	25

SECTION 1

Project Deep-Dive

1. What is PrePro Auto?

PreProAuto is an AI-assisted, human-in-the-loop tabular data preprocessing tool. It combines a statistical preprocessing engine, a local web-based workbench UI, a Jupyter notebook bridge, and optional AI assistance from five LLM providers — all behind a single `pip install`.

The core idea is simple: instead of writing dozens of pandas / sklearn lines to clean a dataset — guessing at imputation strategies, eyeballing outliers, picking scalers from memory — you load your data into the workbench, review AI-generated decision cards for each issue found, approve or override them, and walk away with a clean dataset plus a runnable Python pipeline that reproduces every step.

Key facts at a glance:

Property	Value
Package name	prepro-auto
Install	<code>pip install prepro-auto</code>
Version	1.0.0 (stable)
License	MIT — free for personal and commercial use
Author	Shivanshu Pandey (github.com/Chilliflex)
Python	3.10 / 3.11 / 3.12
Tests	278+ passing (pytest, CI on every push)
Dependencies	FastAPI, pandas, scikit-learn, scipy, reportlab, uvicorn + more
AI (optional)	Groq, OpenAI, Anthropic, Gemini, Mistral — all optional
Engine runs local	Preprocessing engine: 100% local, no telemetry. Optional AI sends column names, dtypes & a few sample values (PII-masked) to your chosen provider only when enabled.

2. Project Goal & Design Philosophy

The goal is to eliminate the most time-consuming and error-prone part of every machine learning project: cleaning the raw data. Data scientists spend a large proportion of project time on preprocessing, and most of that time is spent on decisions that are statistically derivable yet written from scratch each time.

Three design principles guide every decision in the codebase:

Stats beats ML/DL when the answer is derivable. The correct scaler for a Gaussian column is `StandardScaler` — no neural network needs to learn that. **PreProAuto** uses statistical tests (Shapiro-Wilk, KS, IQR, Pearson/Spearman, PSI) where the answer is mathematically derivable, and only involves ML models (Isolation Forest, KNN, MICE) where the data itself must inform the decision.

The human is always in control. Every recommendation is a decision card: the system proposes, the human approves, overrides, or skips. Nothing changes the data without explicit approval. Every change is a versioned snapshot with full undo/redo.

Reproducibility is not optional. Every session produces a runnable Python pipeline script that reproduces the exact cleaning with no **PreProAuto** dependency. The exported script is **transform-only**: all parameters (means, medians, scaler statistics, category maps, target means) are fitted once inside **PreProAuto** and baked into the script as literal values — the script never re-fits on the data it is given. See Section 5 for the full reproducibility guarantee.

3. Target Users

User	How they use PrePro Auto
Data Scientist (ML / DL)	Launch from Jupyter: <code>prepro_auto.launch_file(path)</code> . Clean visually, then <code>session.current()</code> feeds straight into <code>model.fit()</code> . Zero file I/O.
Data Analyst	Run the web UI (<code>prepro_auto</code> in terminal), upload CSV, step through the guided 6-step flow, download clean CSV and audit PDF.
ML Engineer / Data Engineer	Use the exported <code>pipeline.py</code> in Airflow / Prefect / GitHub Actions. Cleaning is automated and reproducible in production.
Student / Researcher	Free, offline, no API key required. A complete preprocessing environment for academic datasets with full explainability of every decision.
Compliance / Audit role	The audit PDF lists every transformation with parameters, before/after stats, and who approved it — a complete provenance trail.

4. System Architecture

PreProAuto has a layered architecture. All layers run locally in a single process when installed via pip — no Docker, no cloud, no network required.

Layer	Technology	Responsibility
Notebook bridge	prepro_auto/__init__.py	quickclean(), launch(), launch_file(), Session(), set_api_key()
REST API	FastAPI + uvicorn	All /api/v1/ endpoints; consumed by both UI and notebook
Preprocessing engine	app/preprocessing/	5 cleaning stages, profiler, transforms, drift
Service layer	app/services/	Orchestrates stages, manages decision cards, export
Privacy layer	app/services/privacy.py	PII detection & masking before any value reaches the AI layer
Model training	app/ml/ + app/services/ ml_service.py	Leakage-safe training: replays the cleaning recipe re-fit per CV fold, ranked model zoo, hyperparameter search, drift check, joblib/ONNX export
Versioning	SQLite + app/db/	Every snapshot is a committed version; undo/redo
Storage	Local filesystem	Raw uploads and snapshots stored in ./storage/
Web UI	app/web/*.html + FastAPI static	6-step guided workbench served at /workbench
AI client	app/services/ llm_client.py	Provider-agnostic LLM wrapper; optional
Export engine	app/services/ export_service.py	CSV/Parquet, audit PDF, transform-only pipeline.py generation

Data flow diagram:

```

User Input (file / DataFrame)
|
[prepro_auto bridge / Web Upload]    (or quickclean() for a headless one-call run)
|
FastAPI REST layer (/api/v1/)
|
Preprocessing Engine
  Profiler (type inference, quality score)
  Missing Values (MCAR/MAR/MNAR + KNN/MICE/median/mode) --> fitted fill values stored
  Outlier Detection (IQR + Hampel + Isolation Forest)
  Scaling (Shapiro-Wilk/KS -> StandardScaler/RobustScaler/Box-Cox...) --> fitted params
    stored
  Correlation + Leakage (Pearson/Spearman, ID detection)
  Encoding (cardinality-routed: label/ordinal/OHE/frequency/target) --> fitted maps stored
|
Decision Cards (human reviews: Approve / Override / Skip)
|
Execute Stage -> Fitted state persisted into TransformationLog.params._fitted
                  -> New Snapshot (versioned, undoable)
|
[Optional AI layer]
  Privacy filter: PII-masked column names + dtypes + small sample (never full dataset)
  --> LLM provider (Groq / OpenAI / Anthropic / Gemini / Mistral)
|
[Optional Step 7: Train]
  Recipe replay: TransformationLog steps with _fitted stripped
  --> sklearn ColumnTransformer, re-fit inside every CV fold (leakage-safe)
  --> ranked model zoo (gradient boosting / random forest / linear / MLP / custom)

```

```
--> best model selected on the chosen metric
--> joblib bundle | scorer script | ONNX
|
Outputs: cleaned DataFrame | CSV/Parquet | Audit PDF | pipeline.py (transform-only) |
        Drift report | trained model bundle
```

5. End-to-End Workflow

Step 1 — Upload / Load




Figure 1: Upload / Load interface.

Drop a CSV, Parquet, Excel, or JSON file onto the upload zone (or call `prepro_auto.launch_file(path)` in a notebook). Encoding and delimiter are automatically detected. The sidebar RAM panel displays live memory availability and upload limits for the current machine.

Example: The `Bengaluru_House_Data.csv` dataset (916 KB) has been successfully loaded. The sidebar shows 2.6 GB free out of 15.3 GB total RAM, along with recommended limits for streaming and memory-intensive operations.

Step 2 — Profile




Figure 2: Profile interface.

The profiler performs automatic dataset analysis by computing semantic data types, missing-value percentages, cardinality statistics, representative sample values, and an overall data-quality score. It **runs automatically** the moment you open the Profile step (no button). This stage is observational only and does not modify the dataset.

Example: Per-column semantic types, missing rates, cardinality statistics, and dataset-level quality metrics are computed in a single profiling pass.

Step 3 — View Dataset




Figure 3: View Dataset interface.

The dataset viewer provides a live tabular representation of the current data state. Users can switch between the original uploaded dataset and the cleaned version at any time. Dataset metadata such as row count, column count, and quality score are updated automatically after every transformation.

Example: The displayed dataset contains 13,320 rows and 9 columns with a quality score of 91.9.

Step 4 — Clean




Figure 4: Clean interface.

The cleaning stage operates through five sequential sub-stages: Missing Values, Outliers, Scaling, Correlation, and Encoding. For each column, the system generates a decision card containing the recommended action, a confidence percentage, a plain-English explanation of the business impact, and a list of alternative strategies. The user reviews each card and selects one of four responses: Approve, Override, Skip, or Drop column. **Each stage runs automatically when you select its tab** — there is no separate “Run stage” button. Executing a stage commits all approved decisions as a new versioned, undoable snapshot *and automatically advances to and runs the next stage*.

Fitted state is persisted at this step. When you execute a stage, the parameters PreProAuto computed — e.g. the median used to fill missing values, the mean and standard deviation the scaler fitted, the category-to-integer map an encoder built, the smoothed (M-estimate) per-category target means — are stored in `TransformationLog.params._fitted`. The exported `pipeline.py` reads these stored values to reproduce the transform without ever re-fitting on new data.

Example: The Encoding stage is active. The `area_type` column has been flagged with an OVER-RIDDEN badge, indicating the user rejected the system’s One-hot encoding recommendation (88% confidence).

Step 5a — Transform: Preset Operations




Figure 5: Transform: Preset Operations interface.

The Preset Operations tab provides 18 built-in, point-and-click data transformations that require no coding. Available operations include: rename, drop, cast (type conversion), fill missing values, filter rows, merge columns, arithmetic operations (+, -, *, /), map values, string transformations (uppercase, lowercase, strip, replace), regex extraction, **group-aggregate (group-by feature engineering)**, sort, deduplicate, and extract number from text. Multiple columns can be selected simultaneously to apply the same operation in a single undoable step.

The **group-aggregate** preset computes, for each row, the aggregate (mean / sum / count / min / max / median / std / nunique) of a value column over its group and broadcasts it back as a new feature column — e.g. *average price per location*.

Step 5b — Transform: Expression (Advanced)




Figure 6: Transform: Expression (Advanced) interface.

The Expression Editor allows users to write pandas expressions against the current DataFrame. The dataframe is exposed as `df`; both `pd` (pandas) and `np` (NumPy) are available. Expressions are validated against a restricted namespace and a denylist of dangerous tokens (imports, file/OS/network access, `.apply`, `.eval`, loops, lambdas, etc.) before execution. Every successful expression application is saved as a new undoable version. See [SECURITY.md](#) for the full mechanism and trust-boundary description.

Trust boundary: The expression sandbox is designed for a **single-user, local-machine** context — the person running the server is the same person whose expressions are evaluated. It is **not** hardened for untrusted multi-tenant input. For multi-user or hosted deployments, run each instance in an isolated container.

Step 5c — Transform: AI Assistant

PrePro Auto
DEVELOPER WORKBENCH

- 1 Upload dataset file
- 2 Profile understand columns
- 3 View data live, before/after
- 4 Clean review & execute
- 5 **Transform** full manual control
- 6 Export data · audit · code
- 7 Train model · leakage-safe

v13 · transform:batch(1)

AI PROVIDER
● Enabled
gpt4o · llama-3.3-70b-versatile
Configure API key

SYSTEM & RAM
3.6 GB free / 15.3 GB total
Upload limit: 8,178 MB · Heavy ops: 408 MB
Refreshes every 15s

STEP 5 OF 6
Transform
Full manual control over the dataset: create, drop, rename, merge, cast, filter, map, and apply math to columns and rows. Every transform is a new version — undo any time.

Preset operations | Expression (advanced) | **AI assistant** | Visualize & metrics | Dashboard (before/after) | Data drift

Chat with your data
Tell it what to change in plain English. It confirms what it understood, shows a 2-row preview, and waits for your approval before touching anything.

remove the text value from BHK column and keep only numeric value in it.

I will remove the text value from the BHK column and keep only the numeric value in it by extracting the number from the column.
This operation assumes that the numeric value in the BHK column is the most relevant information and that the text can be safely removed.

```
extract_number({'column': 'BHK'})
```

BEFORE (FIRST 2 ROWS)

location	BHK	total_sqft	bath	price
Electronic City Phase II	2 BHK	1056	2	39.06999969482422
Chikka Tirupathi	4 Bedroom	2600	5	120

AFTER

location	BHK	total_sqft	bath	price
Electronic City Phase II	2	1056	2	39.06999969482422
Chikka Tirupathi	4	2600	5	120

Approve & apply | Copy to Expression

✓ Applied. Your data is now 13,320 rows × 5 cols. You can keep chatting, or undo from the bar above.

e.g. remove rows where bath is more than 1 above bhk · from the size column keep only the number

Send

Figure 7: Transform: AI Assistant interface.

The AI Assistant tab enables natural-language data transformation. The user describes the desired change in plain English; the assistant interprets the request, generates the corresponding pandas operation, displays the inferred operation code, and presents a two-row before/after preview. No changes are applied until the user explicitly clicks **Approve & apply**.

Privacy: Before sending anything to the LLM provider, **PreProAuto** passes all column names and sample values through `app/services/privacy.py`. Emails, phone numbers, long ID numbers, and card-like patterns are masked to `«email»/«phone»/«id»/«card»`. Columns whose names suggest PII (*email*, *ssn*, *name*, *address*, *phone*, *dob*, etc.) have all their sample values replaced with `«redacted»`. The full dataset is **never** uploaded; only column names, dtypes, and a few masked sample values are sent. Control this with two settings in

app/core/config.py:

LLM_MASK_PII	Default:	Redact PII-looking values before egress on
LLM_SEND_VALUES	Default:	Set False to send no raw values at all on

Step 5d — Transform: Visualize & Metrics




Figure 8: Transform: Visualize & Metrics interface.

The Visualize & Metrics tab provides live, in-workbench charting against the current dataset version. The Chart Builder supports histograms (numeric), bar charts (categorical), and scatter plots. The Metric Counter evaluates user-defined conditions. The Raw vs Current view overlays distributions before and after cleaning. The Ask (AI) option generates a chart from a natural-language description.

Step 5e — Transform: Dashboard (Before and After)




Figure 9: Transform: Dashboard (Before and After) interface.

The Dashboard tab delivers a Power-BI-style summary comparing the raw uploaded dataset to the current cleaned version. Four KPI tiles at the top display row count, column count, missing-value percentage, and duplicate row count, each showing the raw value, the current value, and the net change.

Step 5f — Transform: Data Drift

PrePro Auto
DEVELOPER WORKBENCH

1 Upload
dataset file

2 Profile
understand columns

3 View data
live, before/after

4 Clean
review & execute

5 **Transform**
full manual control

6 Export
data · audit · code

7 Train
model · leakage-safe

Undo Redo

v13 · transform:batch(1)

AI PROVIDER
Enabled
gpg: llama-3.3-70b-verbose

Configure API key

SYSTEM & RAM
3.7 GB free / 15.3 GB total
Upload limit: 8.283 MB · Heavy ops:
414 MB
Refreshes every 15s

Other interfaces:

STEP 5 OF 6

Transform

Full manual control over the dataset: create, drop, rename, merge, cast, filter, map, and apply math to columns and rows. Every transform is a new version — undo any time.

Preset operations Expression (advanced) AI assistant Visualize & metrics Dashboard (before/after) **Data drift**

Compare two datasets for distribution drift
Upload your **reference** dataset (e.g. last month's training data) and a **new** dataset (e.g. recent production data). PrePro Auto computes per-column PSI and KS test for numerics, PSI for categoricals. No model required — works on any two CSV / Parquet / Excel / JSON files.

① Reference dataset
Baseline / training / last period
Choose File Bengaluru_House_Data.csv
Picked: Bengaluru_House_Data.csv

② New dataset
Recent / production / current period
Choose File bhp.csv
Picked: bhp.csv

Detect drift

SIGNIFICANT DRIFT
1 significant · 0 moderate · 4 stable · 5 compared (13,320 ref / 13,200 new)

Schema differs — only in reference: area_type, availability, balcony, society - only in new: bhk, price_per_sqft

Column	Type	PSI	Severity	Detail
total_sqft	categorical	33.6176	significant	20 new categories (99.3% of new rows)
location	categorical	0.008	stable	no new categories
size	categorical	0.0001	stable	no new categories
bath	numeric	0	stable	ref mean 2.6926 → new mean 2.6911, KS p=1
price	numeric	0	stable	ref mean 112.5656 → new mean 112.2762, KS p=1

PSI < 0.10 stable · 0.10–0.25 moderate · ≥ 0.25 significant (industry-standard threshold). Drift means a model trained on the reference data may behave differently on the new data — investigate or retrain.

Figure 10: Transform: Data Drift Detection interface.

The Data Drift tab allows users to upload a second dataset and compare its column-level distributions against the reference dataset. The engine computes the Population Stability Index (PSI) for both numeric and categorical columns and the Kolmogorov-Smirnov (KS) two-sample statistic for numeric columns. Results are classified into three severity bands: stable (PSI < 0.10), moderate (PSI 0.10–0.25), and significant (PSI ≥ 0.25).

Step 6 — Results & Export

PrePro Auto
DEVELOPER WORKBENCH

- 1 Upload
dataset file
- 2 Profile
understand columns
- 3 View data
live, before/after
- 4 Clean
review & execute
- 5 Transform
full manual control
- 6 Export
data • audit • code
- 7 Train
model • leakage • safe

Undo Redo

v13 • transform:batch(1)

AI PROVIDER
Enabled
gpus: 8x NVIDIA 3.3 70b-variant

Configure API key

SYSTEM & RAM
3.6 GB free / 15.3 GB total
Upload limit: 8,143 MB • Heavy ops: 408 MB
Refreshes every 15s

Other interfaces:
Swagger /docs
Review page /review
AI settings (API key)

job:
6bc530a7-8aaf-4a76-a774-8919708ba492

Results & export

See exactly what preprocessing changed — the quality score, the column types before vs after, and a data sample — then download your deliverables.

Before & after
Comparison of the raw upload against the cleaned result. Refresh

BEFORE • RAW

91.9

13,328 rows • 9 cols

AFTER • CLEANED • V13

96

13,328 rows • 5 cols

→ +4.1

Column types: before → after

COLUMN	BEFORE TYPE	AFTER TYPE
area_type	categorical_nominal	+ (dropped)
availability	categorical_nominal	+ (dropped)
location	unknown	+ unknown
size	categorical_nominal	+ (dropped)
society	unknown	+ (dropped)
total_sqft	unknown	+ unknown
bath	numeric_discrete	+ numeric_discrete
balcony	numeric_discrete	+ (dropped)
price	numeric_continuous	+ numeric_continuous
BHK	(new)	+ numeric_discrete

Data sample: before

area_type	availability	location	size	society	total_sqft	bath	balcony	price
Super built-up Area	19-Dec	Electronic City Phase II	2 BHK	Coomee	1856	2	1	39.87
Plot Area	Ready To Move	Chikka Tirupathi	4 Bedroom	Theanmp	2688	5	3	128
Built-up Area	Ready To Move	Uttarahalli	3 BHK		1448	2	3	62
Super built-up Area	Ready To Move	Lingadheeranahalli	3 BHK	Solowre	1521	3	1	95
Super built-up Area	Ready To Move	Kothanur	2 BHK		1288	2	1	51

Data sample: after

location	BHK	total_sqft	bath	price
Electronic City Phase II	2	1856	2	39.87
Chikka Tirupathi	4	2688	5	128
Uttarahalli	3	1448	2	62
Lingadheeranahalli	3	1521	3	95
Kothanur	2	1288	2	51

Download
Take your results away.

Clean data
The processed dataset as CSV or Parquet.

CSV ▼ Download

Audit report
A PDF listing every transformation and who decided it.

Download PDF

Pipeline code
A standalone Python script reproducing the cleaning.

Download .py

Figure 11: Results & Export interface.

The Results & Export step provides a final summary of all preprocessing activity and three downloadable artefacts: the cleaned dataset (CSV or Parquet), a full Audit PDF, and a standalone pipeline.py script.

Reproducibility guarantee — no train/serve skew, no leakage.

The exported pipeline.py is **transform-only**. Every parameter each step needs was **fitted once inside PreProAuto** and is baked into the script as a literal value. The script never re-fits on the data it is given. Consequences:

- Running on a held-out / test set applies the *training* statistics — no test-set leakage, no train/serve skew.

- One-hot encoding always emits the **same columns in the same order**, regardless of which categories appear in new data (missing → 0, unseen → `__other__`).
- Target encoding uses M-estimate **smoothing** — rare-category means are shrunk toward the global mean, curbing overfitting and target leakage — and reuses the stored map (no recompute, no train/serve skew on new data).
- KNN/MICE: the pipeline reuses the **median learned at clean time**. (Full KNN/MICE cannot be deterministically re-applied to arbitrary new rows without skew; pickle-sidecar export is on the roadmap.)
- Input format is auto-detected: CSV, Parquet, Excel (.xlsx/.xls), and JSON are all handled — not just CSV.

Verified by `tests/test_reproducibility.py` and an end-to-end run on held-out data. Run: `python pipeline.py raw.csv ready.csv`

Example: The dataset quality score improved from 91.9 (raw, 13,320 rows, 9 columns) to 97.9 (cleaned, 13,304 rows, 5 columns), a gain of +6 points.

Step 7 — Train a Model




Figure 12: Train — Setup interface.

Once a dataset is clean, Step 7 — reached from the left sidebar (after Export) or directly at `/train?job=<job_id>` — fits and compares candidate models on top of the cleaning recipe, without leaving the browser. Two fields drive everything else: the **target column** and a one-line **business goal** in plain English (“estimate house price” vs. “catch fraud” vs. “score/rank leads”) — the goal selects the primary metric (RMSE, PR-AUC, ROC-AUC, macro-F1) automatically, so accuracy is never the silent default on imbalanced data.

Leakage-safe by construction — not just leakage-checked.

Training does **not** reuse the all-rows fitted values from Step 4 (`TransformationLog.params._fitted`).

Instead the cleaning recipe is replayed with those fitted values **stripped**, compiled into a scikit-learn `ColumnTransformer`, and re-fit **inside every cross-validation fold** on that fold's training rows only. Consequences:

- Target encoding (M-estimate smoothed) cannot leak test-fold labels into training features — it is re-fit per fold, not fit once on all rows.
- Numeric-looking string columns (e.g. "1,843") are coerced and imputed rather than silently dropped; any categorical left unencoded by the recipe is still encoded — the training matrix is guaranteed fully numeric before it reaches a model.
- Reported cross-validated scores match what you would actually see in production, not an inflated all-rows-fit number.

Verified by `tests/test_ml_leakage.py` and `tests/test_ml_recipe.py`.




Figure 13: Train — ranked leaderboard after training.

Suggest, then approve. Clicking **Suggest models** ranks the candidate models most-relevant-first for the detected problem type and dataset size, flags one as *recommended*, and fills in the metric — **before** anything trains. Each candidate is shown as a card with its family (linear, ensemble, boosting, neural) and a one-line reason. Pick one and **Approve & train selected**, or skip straight to **Train all & compare**.

Built-in model zoo: gradient boosting (LightGBM/XGBoost/CatBoost if installed, else scikit-learn's HistGradientBoosting), random forest, a linear/logistic baseline, and a small neural network (MLP) — included on request, though it rarely beats boosting on tabular data and is shown last, not recommended by default.

Bring your own model. A separate **Custom model** card accepts any scikit-learn-compatible estimator by import path (e.g. `sklearn.svm.SVR`, `sklearn.linear_model.ElasticNet`, `xgboost.XGBRegressor`), a family label (ML / statistical / neural / ensemble), and optional JSON init params. Only classes from an allowlisted set of libraries (scikit-learn, XGBoost, LightGBM, CatBoost) can be loaded this way; anything else is rejected and surfaces as a single failed leaderboard row rather than executing.

Hyperparameter tuning is opt-in via a dropdown: *None* (defaults, fastest), *Randomized search*, *Grid search*, *Successive halving*, or *Optuna / Bayesian* (if installed). Tuned parameters are shown per model on the leaderboard alongside the cross-validated and holdout scores and the fit time.




Figure 14: Train — live test playground.

Once a model is trained, the **test playground** builds an input form straight from the model's feature schema — enter values, click **Predict**, and see the live prediction (and class probabilities, for classification) without writing any code.




Figure 15: Train — export options.

Export the winning model as a **joblib bundle** (preprocessing + model in one object, ready to load in Python), a standalone **scorer script** carrying the feature schema, or **ONNX** for

cross-language serving. A built-in **drift check** (PSI/KS, reusing the same drift engine as Step 5f) compares new inference-time data against the training distribution.

PrePro Auto
DEVELOPER WORKBENCH

1 Upload
dataset file

2 Profile
understand columns

3 View data
live, before/after

4 Clean
review & execute

5 Transform
full manual control

6 Export
data · audit · code

7 Train
model · leakage-safe

Undo Redo

v0 · transform:batch(1)

AI PROVIDER
Enabled
groq · llama-3.3-70b-versatile

STEP 7 OF 7

Train a model

Train a model on top of your cleaning recipe. Preprocessing is re-fit inside every cross-validation fold from the raw data, so the scores are leakage-free and match what you'll see at serving time — not the inflated numbers you get from fitting on all rows first.

- 1 · **Setup** Name the target and your goal in plain English — the goal picks the metric ("catch fraud" → PR-AUC, "estimate price" → RMSE).
- 2 · **Suggest** See the candidate models ranked for your data, one **flagged** as recommended. Approve it, pick another, or add your own.
- 3 · **Train** Compare on a held-out split, read the leaderboard, test predictions live, and export the pipeline.

Setup

What you're predicting and why. Everything else can stay on defaults.

TARGET COLUMN: price

BUSINESS GOAL: estimate the price based on columns input

The column you want to predict.

Steers the metric. Plain English is fine.

Validation & search

TEST SIZE: 0.2

CV FOLDS: 5

HYPERPARAMETER TUNING: None (defaults, fastest)

LIMIT TO MODELS (OPTIONAL)

blank = all; or e.g. gradient_boosting, random_forest

Comma-separated ids. Leave blank to compare everything, including custom models below.

Suggest models Train all & compare →

Figure 16: Train — selecting the data source for a run.

Two ways to train, matching the two ways to clean: from the **web workbench** as described above (*strict* guarantee — always trains from raw data + recipe), or from a **notebook** via `session.train(...)` / `result.train(...)` / the top-level `prepro_auto.train(data, target, ...)` (see Section 10) — *strict* when called on a `Session` or `CleanResult` object, *trusted-input* when called on a bare `DataFrame` or file path (the data is scanned for obvious leaks, but there is no recipe to safely replay).

Example: On the Bengaluru house-price dataset, gradient boosting was recommended and trained to holdout $R^2 \approx 0.64$ with all 8 raw feature columns used (numeric-looking strings such as `total_sqft` coerced rather than dropped).

6. How It Works (Methods at a Glance)

PreProAuto uses only field-standard, transparent methods — no black boxes. Where the right answer is derivable from a column's distribution it uses a statistical test; machine learning is used only where the data itself must inform the choice. No deep learning is used.

Stage	Methods used
Profiling & typing	Regex / match-rate type inference; dictionary + heuristic semantics (optional AI hint).
Missing values	Mechanism check (MCAR / MAR / MNAR), then median, mode, KNN, MICE, or a was-missing flag. <i>Fitted fill values are persisted for transform-only export.</i>
Outliers	IQR (Tukey) + modified z-score (Hampel / MAD); Isolation Forest for multivariate rows.
Scaling	Normality test (Shapiro-Wilk / KS), then Standard, Robust, MinMax, Box-Cox, Yeo-Johnson, or log1p. <i>Fitted scaler params (mean, std, center, lambda...) are persisted.</i>
Correlation & leakage	Pearson / Spearman redundancy; ID-uniqueness and target-leak detection.
Encoding	Cardinality-routed: label, ordinal, one-hot, frequency, or smoothed target encoding. <i>Fitted maps, frozen category sets, and M-estimate smoothed target means are persisted.</i>
Drift (monitoring)	Population Stability Index (PSI) + Kolmogorov-Smirnov two-sample test.

Every recommendation is shown as a decision card with a confidence score and a plain-English reason, so nothing is applied without your understanding and approval.

7. Dataset Size Limits

PreProAuto sets limits based on the machine's available RAM at request time — not a fixed number. The reason is that pandas loads the entire DataFrame into memory, expanding a raw CSV file by approximately 2–10× (typically 6× for mixed-type data).

Limit type	Formula	Typical value (8 GB laptop)	Hard ceiling
Streamable (profiling, scaling, encoding, IQR outliers, filters)	$\min(10 \text{ GB}, \max(200 \text{ MB}, \text{available_RAM} * 0.55 * 4))$	~1.8 GB raw file	10 GB
Heavy in-memory (KNN, MICE, Isolation Forest)	$\max(100 \text{ MB}, \text{available_RAM} * 0.55 / 5)$	~180 MB raw file	10 GB

Practical guidance: For datasets under 200 MB, all features including KNN/MICE and Isolation Forest work on any modern laptop. For 200 MB–2 GB, profiling, scaling, encoding, and export work fine; KNN/MICE may be unavailable — use median/mode instead. Above 2 GB, use Dask or Spark for the heavy operations.

SECTION 2

Complete User Guide — Notebook & Web UI

8. Installation

Base install (offline — no AI features, but fully functional):

```
pip install prepro-auto
```

With a specific AI provider:

```
pip install prepro-auto[groq]      # Groq -- free tier, fastest
pip install prepro-auto[openai]   # OpenAI / GPT-4
pip install prepro-auto[anthropic] # Anthropic Claude
pip install prepro-auto[gemini]   # Google Gemini
pip install prepro-auto[mistral]  # Mistral
pip install prepro-auto[ai]       # All five providers at once
```

Important for Jupyter users: after installing in a notebook cell with `%pip install prepro-auto`, restart the kernel before importing. The `prepro_auto` command in a Jupyter cell prints the module object — it does **NOT** start the server. See the *Web UI — Step-by-Step Guide* section for the correct way to start the web UI; inside a notebook use `prepro_auto.quickclean(...)`, `launch_file(...)` or `launch(df)` instead.

9. Data Privacy

What leaves your machine?

The preprocessing engine is 100% local. Profiling, cleaning, scaling, encoding, drift, export, versioning, and storage all run in-process with no network calls and no telemetry.

The **optional AI features are the only part that uses the network**. When you enable a provider, **PreProAuto** sends **only**: column names and dtypes, a few sample values per column (default 3–4, never the full dataset), and a 2-row before/after preview when the AI assistant applies a transform.

PII masking: Before any sample value leaves the machine, `app/services/privacy.py` redacts: emails (→ «email»), phone numbers (→ «phone»), long ID-like numbers (→ «id»), card-like patterns (→ «card»). Columns whose names suggest PII (email, ssn, name, address, phone, dob, account, patient, ...) have **all** sample values replaced with «redacted».

Two config controls in `app/core/config.py`:

LLM_MASK_PII	Default:	Redact PII-looking values before egress
	True	
LLM_SEND_VALUES	Default:	Set False: send no raw values at all
	True	

Will the provider “learn” my data? Major providers generally do not train on API traffic; several offer zero-retention modes. Check your provider’s data-processing terms. **For zero network egress**, leave AI disabled, or point it at a local OpenAI-compatible endpoint (e.g. Ollama).

10. Notebook Interface — All Functions

All notebook functions live in the `prepro_auto` module. Import once and use throughout your notebook.

```
prepro_auto.quickclean(data, *, target=None, threshold=0.90, apply_all=False, stages=None, verbose=True)
```

Returns: `CleanResult` (with `.df`, `.report`, `.job_id`)

The fastest path — one call, no browser. Runs the entire engine headlessly and applies every decision at or above `threshold` confidence. Pass `apply_all=True` for full auto-pilot. `data` may be a pandas `DataFrame` or a path to a CSV/Excel/JSON/Parquet file.

```
import prepro_auto

result = prepro_auto.quickclean("sales.csv", target="churn")
df = result.df          # cleaned, model-ready DataFrame
print(result)          # per-stage report: applied vs. left for review

X = df.drop(columns=["churn"]); y = df["churn"]
model.fit(X, y)

# Full auto-pilot:
result = prepro_auto.quickclean("sales.csv", target="churn", apply_all=True)
```

```
prepro_auto.launch_file(file_path, *, domain='general', port=None, open_browser=False)
```

Returns: `Session`

Recommended starting point. Reads a file from disk with auto-encoding and auto-delimiter detection. Handles UTF-8, Latin-1, cp1252, BOM-prefixed files silently.

```
session = prepro_auto.launch_file(r'C:\data\sales.csv')
# Click: http://127.0.0.1:8721/workbench?job=...
```

```
prepro_auto.launch(df, *, domain='general', port=None, open_browser=False)
```

Returns: `Session`

Pass an existing pandas `DataFrame` directly — no file I/O. The original `df` is never modified; `session.current()` returns the cleaned version.

```
import pandas as pd, prepro_auto
df = pd.read_csv('data.csv')
session = prepro_auto.launch(df)
cleaned = session.current()
```

```
prepro_auto.Session(job_id, port)
```

Returns: Session

Reconnect to an existing session by ID and port.

```
session = prepro_auto.Session('64c15bf6-...', port=8721)
```

```
prepro_auto.set_api_key(provider, api_key, *, model=None)
```

Returns: dict with ok/verified/provider/model/reason

Configure an AI provider for the current session. Held in memory only — never written to disk.

```
result = prepro_auto.set_api_key('groq', 'gsk_...')
# result: {ok: True, verified: True, provider: 'groq', model: 'llama-3.3-70b-
          versatile'}
```

Session object methods

```
session.current()
```

Returns: pandas DataFrame

Returns the current active-version DataFrame. Thread-safe.

```
X = session.current().drop(columns=['target'])
y = session.current()['target']
model.fit(X, y)
```

```
session.update(df)
```

Returns: the same Session

Push a notebook-edited DataFrame back to the workbench as a new undoable version.

```
import numpy as np
df2 = session.current()
df2['log_price'] = np.log1p(df2['price'])
session.update(df2)
```

```
session.url / session.job_id / session.port
```

```
print(session.url)      # http://127.0.0.1:8721/workbench?job=64c15bf6-...
print(session.job_id)   # UUID for direct REST API calls
print(session.port)     # e.g. 8721
```

```
prepro_auto.train(data, target, *, business_goal=None, models=None, tuning='none',
cv_folds=5, explain=False)
```

Returns: TrainResult (with .best, .leaderboard, .leakage, .predict(...), .export(path))

Train a model on top of the cleaning recipe without opening a browser. `data` may be a `Session`, a `CleanResult`, a bare `DataFrame`, or a file path. From a session or `CleanResult`, training is **strict**: it re-derives the preprocessing from the raw data and cleaning recipe, re-fit inside every cross-validation fold (see the leakage-safe guarantee in Section 5). From a bare `DataFrame` or path, training is **trusted-input**: there is no recipe to safely replay, so the data is scanned for obvious leaks and used as given.

```
import prepro_auto

# Strict / leakage-safe -- straight from a clean() session:
result = prepro_auto.quickclean("houses.csv", target="price", apply_all=True)
trained = result.train(target="price", business_goal="estimate price")
print(trained)                                # leaderboard + metric + leakage guarantee
trained.predict({"location": "Whitefield", "total_sqft": 1200, "bath": 2})
trained.export("model.joblib")

# Also callable on a live workbench session:
session = prepro_auto.launch_file("houses.csv")
trained = session.train(target="price", business_goal="estimate price", tuning="
    random")

# Trusted-input -- a bare DataFrame or file path you already trust:
trained = prepro_auto.train(my_clean_df, target="price")
```

Beginner to Advanced — Notebook Progression

Level	What you do	Code pattern
Fastest	One headless call — no browser, no clicking	<code>result = prepro_auto.quickclean(path, target='y')</code> <code>df = result.df</code>
Beginner	Load a file, open the workbench, clean visually, get the clean <code>DataFrame</code>	<code>session = prepro_auto.launch_file(path)</code> <code>cleaned = session.current()</code>
Intermediate	Enable AI, use <code>set_api_key</code> , push notebook edits back to the UI	<code>prepro_auto.set_api_key('groq', key)</code> <code>session.update(my_edited_df)</code>
Advanced	Reconnect to an existing session, make REST API calls directly	<code>session = prepro_auto.Session(job_id, port)</code> <code>import requests;</code> <code>requests.get(f'{session.url}/api/v1/...')</code>

11. Web UI — Step-by-Step Guide

Starting the web UI: Open Command Prompt (not Jupyter) and run:

```
prepro_auto          # default: http://127.0.0.1:8000
prepro_auto --port 9000  # custom port
prepro_auto --host 0.0.0.0 # accessible on local network
```

Then open <http://127.0.0.1:8000/workbench> in your browser.

Step 1 — Upload

- Drag-and-drop a CSV, Parquet, Excel, or JSON file onto the upload zone.
- Encoding and delimiter are auto-detected — no configuration needed.

Step 2 — Profile

- **Profiling runs automatically** when you open the Profile step — there is no button to click.
- You see: semantic type, missing rate, unique count, sample values, quality score (0–100).
- No changes are made at this stage — it is observation only.

Step 3 — View data

- Toggle between **Original (raw)** and **Current** (after any cleaning) views.
- Use **Refresh** to pull the latest state after making changes in other steps.

Step 4 — Clean (HITL)

- Five stage tabs: **Missing Values** | **Outliers** | **Scaling** | **Correlation** | **Encoding**.
- **Selecting a stage tab runs it automatically** — there is no ‘Run stage’ button. Opening the Clean step auto-runs **Missing Values** first.
- **Quick clean** (one button) runs *all five* stages and applies every recommended decision in a single click — fully undoable. It is the visual equivalent of `prepro_auto.quickclean(..., apply_all=True)`.
- For each card: **Approve**, **Override**, **Skip**, or **Drop column** (or **Approve all pending**).
- Click **‘Execute stage’** to commit — this **auto-advances to and runs the next stage**, refreshes the issue-count badges, and **persists fitted state** for the reproducible export.
- **Undo / Redo** buttons are always available in the sidebar.

Step 5 — Transform

- 18 preset operations: rename, drop, cast, fillna, filter rows, merge columns, math (+/-/*//), map values, string ops, regex extract, **group-aggregate**, sort, deduplicate, extract number.
- Expression editor: restricted pandas expressions evaluated in a denylist sandbox.
- AI assist (requires API key + PII masking active): describe in English, AI proposes, you confirm.

Step 6 — Export

- Download cleaned data as CSV or Parquet.
- Download audit PDF: every decision documented with parameters and before/after stats.
- Download pipeline.py: **transform-only**, format-aware (CSV/Parquet/Excel/JSON), no **PreProAuto** dependency. See reproducibility guarantee above.

Step 7 — Train

- Set the target column and a one-line business goal; click **Suggest models** to see ranked candidates with one flagged as recommended, or skip straight to **Train all**.
- Optionally add a **custom model** by import path (scikit-learn / XGBoost / LightGBM / CatBoost only), and pick a **hyperparameter tuning** strategy.
- Preprocessing is re-fit inside every cross-validation fold from the raw data — leakage-safe by construction, not just leakage-checked (see Section 5).

- Test the winning model live in the **playground**, then export as joblib, a scorer script, or ONNX.

12. All Outputs Explained

Output	How to get it	What's in it	Where to use it
CleanResult (headless)	<code>prepro_auto.quickclean (data)</code>	<code>.df</code> , <code>.report</code> , <code>.job_id</code>	One-call cleaning; <code>result.df</code> feeds directly into <code>model.fit(X, y)</code> .
Cleaned DataFrame	<code>session.current()</code> in notebook	Active-version pandas DataFrame	Feed directly into <code>model.fit(X, y)</code> without saving intermediate files.
Cleaned CSV	Export step or GET <code>.../export/data?format=csv</code>	UTF-8 CSV	Share with teammates, BI tools (Tableau, Power BI), and data repositories.
Cleaned Parquet	Export step or GET <code>...?format=parquet</code>	Column-oriented binary	Large datasets, feature stores, and distributed Spark/Dask pipelines.
Audit PDF	Export step or GET <code>.../export/audit</code>	Every transformation, parameters, before/after stats, approver, quality delta	Compliance documentation, model cards, audits, and governance handoffs.
Pipeline script (pipeline.py)	Export step or GET <code>.../export/pipeline</code>	Transform-only script using stored fitted params. Format-aware input. No PreProAuto dependency.	Production pipelines, Airflow/Prefect workflows, CI/CD automation, and batch preprocessing jobs.
Drift report	POST <code>/api/v1/drift/compare</code>	Per-column PSI + KS verdict with severity bands	Model monitoring and automated alerts when <code>significant_drift</code> is detected.
Trained model bundle	Step 7 Export, or GET <code>.../ml/runs/{run_id}/export?fmt=</code>	joblib (preprocessing + model), a standalone scorer <code>.py</code> , or ONNX	Serve predictions in Python without PreProAuto , or in another language via ONNX runtimes.

13. AI Providers Setup

AI features are entirely optional. The full preprocessing engine runs offline without any API key. AI adds: plain-English transform proposals, column advice, and multi-turn conversation about your data. All AI egress is PII-masked (see Section 9).

Provider	Best for	Install	Get key	Free?
Groq	Fast, cheap, testing	<code>pip install prepro-auto[groq]</code>	console.groq.com	Yes
OpenAI	GPT-4o, strongest quality	<code>pip install prepro-auto[openai]</code>	platform.openai.com	No
Anthropic	Claude — reasoning	<code>pip install prepro-auto[anthropic]</code>	console.anthropic.com	No
Gemini	Google integration	<code>pip install prepro-auto[gemini]</code>	aistudio.google.com	Yes
Mistral	European / OpenAI-compatible	<code>pip install prepro-auto[mistral]</code>	console.mistral.ai	No

Three ways to configure a key:

```
# 1. Notebook (safest -- stays in memory only, never written to disk)
prepro_auto.set_api_key('openai', 'sk-...')

# 2. Web UI -- click 'Configure API key' in the sidebar, paste key, click Test &
    apply

# 3. .env file (survives restarts -- never use on shared or hosted machines)
LLM_PROVIDER=groq
GROQ_API_KEY=gsk...
```

14. Troubleshooting

Problem: `prepro_auto` typed in Jupyter prints the module object

Solution: Use `prepro_auto.launch_file(path)` or `prepro_auto.launch(df)` in a notebook. The CLI command only works in a terminal.

Problem: `ERR_CONNECTION_REFUSED` at <http://127.0.0.1:8000>

Solution: The server is not running. Open Command Prompt and run `prepro_auto`. Wait for the 'Starting PrePro Auto workbench at...' message before opening the URL.

Problem: `AttributeError: module 'prepro_auto' has no attribute 'launch_file'`

Solution: Old version installed. Run: `pip install --upgrade prepro-auto`

Problem: `UnicodeDecodeError` on CSV load

Solution: Use `prepro_auto.launch_file(path)` instead of `pd.read_csv(path)`. `launch_file` auto-detects encoding (UTF-8, Latin-1, cp1252, BOM) silently.

Problem: KNN or MICE imputation fails on large dataset

Solution: Check `GET /api/v1/system/limits` to see `heavy_max_mb` for your machine. Use median/mode imputation instead, or run on a machine with more available RAM.

Problem: Pipeline export uses median instead of re-running KNN/MICE on new data

Solution: This is intentional and correct. KNN/MICE cannot be deterministically re-applied to arbitrary new rows without re-fitting on (and therefore leaking) the new data. The exported pipeline reuses the **median learned at clean time**. This preserves the train/serve contract. Pickle-sidecar export (which would allow the full imputer to be reused) is on the roadmap.

Problem: Drift detection returns ‘no common columns’

Solution: Ensure both reference and new datasets have the same column names.

Questions, bugs, or feature requests?

Open an issue at https://github.com/Chilliflex/prepro_auto

PyPI page: <https://pypi.org/project/prepro-auto>

Security disclosures: see SECURITY.md in the repository.